

RICK ASHFORD

Director, Sales Engineering

Pflugerville, Texas • rick@theashfords.org • [linkedin.com/in/rickashford](https://www.linkedin.com/in/rickashford)

PROFESSIONAL SUMMARY

Sales engineering leader with 17+ years building and leading high-performance SE organizations across North America and Latin America, inside a 27-year career that started in code. Grew strategic partner bookings from single-digit to double-digit millions and sustained employee eNPS above 90. Deep expertise in enterprise infrastructure, Kubernetes, and container security — extended into AI platform work, from enabling partner AI practices on SUSE's LLM and RAG stack to running local inference, RAG, and autonomous agents on a self-built cluster. Bilingual in English and Portuguese.

CORE COMPETENCIES

Technical Pre-Sales Leadership	Team Development & Player-Coach Model	MEDDPICC Sales Methodology
Revenue Growth & Partner Ecosystems	AI Platforms, Cloud & Kubernetes	Cross-Functional GTM Strategy
Solution Architecture & POC Design	Talent Recruiting & SE Mentorship	Bilingual (English / Portuguese)

PROFESSIONAL EXPERIENCE

Director of Solution Architects, North American Ecosystem • SUSE • Nov 2022 – Nov 2025 •

Remote

- Drove significant revenue growth across top strategic partners over two years: CDW grew from \$8M to \$17M, WWT from \$2M to \$4.5M, and Presidio from ~\$100K to \$2M — achieved through deepened executive relationships, aligned technical resources, and jointly developed go-to-market motions.
- Enabled partners to build their own AI practices on SUSE's Kubernetes-based AI platform — an LLM and RAG stack built on Ollama, Open WebUI, and Milvus — translating an emerging product line into deliverable partner capability across the ecosystem.
- Designed and deployed a standardized partner demo framework adopted across all US channel and reseller partners — enabling lower-proficiency partners to deliver on-message demonstrations and accelerating time-to-demo for advanced partners.
- Recruited, mentored, and developed a solution architect to SUSE SE of the Year within three years — a direct result of deliberate coaching investment and structured performance development.
- Led technical enablement for three acquired product lines (Rancher, NeuVector, StackState), designing partner-facing readiness content that translated new acquisitions into revenue-producing ecosystem capabilities across the channel.
- Defined and executed the North American ecosystem go-to-market technical strategy across Channel & Distribution Partners, Cloud and Managed Service Providers, and Global Systems Integrators — serving as senior technical voice and escalation point for complex partner integrations.
- Partnered with Product and Engineering to translate field learnings and ecosystem feedback into actionable roadmap input, influencing platform direction based on real-world deployment challenges across the partner network.

Sales Engineering Manager, Western/Central US & LATAM • SUSE • May 2016 – Nov 2022 •

Round Rock, TX

- Built and led a team of 18 sales engineers across North America (10) and LATAM (8) — recruiting and developing talent that produced four SUSE SE of the Year honorees over the tenure.
- Hit 100% of quota or higher every year across the region, driven by disciplined MEDDPICC-based pipeline qualification, rigorous forecast management, and consistent technical execution at the field level.

- Led the technical pre-sales motion for a then-company-record \$20M, three-year enterprise deal — personally orchestrating solution design, proof-of-concept delivery, and executive alignment from first engagement through close.
- Co-created and delivered Expert Days, SUSE's highest-performing global marketing event series — presenting at 10 North American events and contributing to 20 additional sessions worldwide, reaching audiences of 20 to 100 technical and business decision-makers per event.
- Sustained employee eNPS above 90 across the region, reflecting a team culture built on high standards paired with genuine investment in people — and retained the team through it: four departures across nine years of management, only one for performance.
- Partnered with a major aerospace company to develop a STIG and auditing schema enabling classified workloads to run on SUSE Linux — a strategic win that expanded SUSE's footprint in defense and regulated industries.
- Qualified for President's Club as a sales engineering manager, recognized for both personal contribution and regional team performance against aggressive targets.
- Represented SUSE as a speaker at SUSECon five times, covering BTRFS, image building with Kiwi, and PaaS with Cloud Foundry — reinforcing the player-coach model by staying technically active and field-credible alongside the team.

Technical Sales Specialist • SUSE • Aug 2008 – May 2016 • Austin, TX

- Twice achieved then-company-record deal sizes as an individual contributor — demonstrating consistent ability to win complex, high-value enterprise opportunities in competitive environments.
- Qualified for President's Club four times as an individual contributor across an eight-year tenure, reflecting sustained top-quartile performance and persistent pipeline contribution.
- Represented SUSE at the OpenStack Summit three times, establishing technical credibility in the open infrastructure community and expanding SUSE's visibility among enterprise cloud decision-makers.
- Delivered hands-on pre-sales expertise for enterprise solutions including SUSE Linux, OpenStack, and emerging container technologies — conducting live demos, proofs-of-concept, and solution workshops for Fortune 500 clients.
- Built trusted, long-term customer relationships through consistent technical excellence and solution advocacy, directly contributing to record-setting pipeline and market share growth.

Deployment Engineer / Sales Engineer → Manager of Operations • Catalis, Inc. • Jan 2004 – Dec 2007 • Austin, TX

Healthcare IT startup delivering electronic medical records (EMR) systems to clinical practices. Fluid startup roles — deployment engineering and sales engineering were the same job.

- Rescued a deal the company could not deliver: a rep sold an ASP (hosted) model Catalis had not planned to build for two more years. Built the VM hosting environment that made it real — \$200K+ in new sales within three months — then built the sales engineering function from nothing so a rep could never again sell a product that did not exist.
- Automated server build procedures, reducing build time from 8–12 hours to 1–2 hours and cutting customer-impacting issues by 95% — a foundational infrastructure improvement that scaled delivery capacity without additional headcount.
- Standardized clinic implementation through checklists, documentation, and post-mortem reviews — cutting average on-site installation from 2–3 days to 4 hours.
- As Manager of Operations (final months), oversaw IT, network operations, training, support, and account management — saving four at-risk clinic accounts worth \$300K in immediate revenue and ~\$5K/month recurring.

SELF-DIRECTED TECHNICAL WORK

NemoClaw Ops Agent • github.com/darthzen/nemoclaw-ops-agent

- Autonomous delivery and runtime operations agent running on a live Kubernetes cluster — detect, diagnose, escalate to Slack, act on approval, audit. Runs entirely on on-premises inference with a documented v1 → v3 autonomy ladder. Security model: two scoped ServiceAccounts, a repo-scoped fine-grained token, and a default-deny NetworkPolicy proven to hold against a root/SYS_ADMIN container.

ash4d Home Lab / lab-fleet • github.com/darthzen/lab-fleet

- Single-node k3s cluster running a production-shaped AI stack: local LLM inference (Ollama/Qwen3) pinned to a Tesla V100 by GPU UUID, image generation isolated on a second GPU, and RAG over Milvus with a FastAPI

retrieval service and a CronJob that incrementally re-indexes and embeds documents. Published as install-ordered config-as-code and managed with Rancher Fleet GitOps.

ollama-code-mcp • github.com/darthzen/ollama-code-mcp

- Model Context Protocol server that delegates coding tasks from a cloud agent to a local Ollama model — keeping inference, and the code it touches, on private infrastructure.

EDUCATION

Master of Information Technology — American InterContinental University, 2003

Bachelor of Arts, Near Eastern Studies — Brigham Young University, 2001

TECHNICAL SKILLS

AI & ML Infrastructure:	Local LLM inference (Ollama — Qwen3, Llama 3.1, DeepSeek-Coder); RAG pipelines (Milvus vector DB, FastAPI retrieval, nomic-embed-text embeddings); GPU scheduling on Kubernetes (NVIDIA device plugin, DCGM, UUID pinning); ComfyUI
Agentic Systems:	Model Context Protocol (MCP) — authored servers, tool whitelisting, OpenAPI gateway (mcpo); OpenClaw / NemoClaw / OpenShell; Kubernetes agent sandboxing (agents.x-k8s.io); Open WebUI; approval-gated autonomy design
AI Security & Governance:	Zero-egress on-premises inference; NetworkPolicy egress isolation; least-privilege RBAC and scoped ServiceAccounts; Pod Security Admission; auditable action contracts
Containers & Orchestration:	Kubernetes, k3s, Rancher, Rancher Fleet (GitOps), Harvester, Longhorn, MetalLB, Traefik, cert-manager, Harbor, Podman, Docker
Container & Infra Security:	NeuVector container security, Kubernetes RBAC, STIG hardening and audit schema design, classified/defense workload enablement
Identity & Directory:	Novell Open Enterprise Server / eDirectory (co-sell), LDAP directory services (Active Directory alternative shipped in SUSE Linux Enterprise), Kubernetes RBAC, certificate/PKI-backed auth
Cloud Platforms:	AWS, Azure, GCP
Enterprise Infrastructure:	SUSE Linux Enterprise, OpenStack, HPC, Storage, ARM/aarch64, VMware
Automation & Observability:	Python, Bash, Ansible, Salt, Chef, Puppet; Git, GitHub Actions, Jenkins, CI/CD pipelines; Prometheus, Alertmanager, DCGM
Sales Tools:	Salesforce, Gong, MEDDPICC methodology
Languages:	English (Native), Portuguese (Full Professional Proficiency)